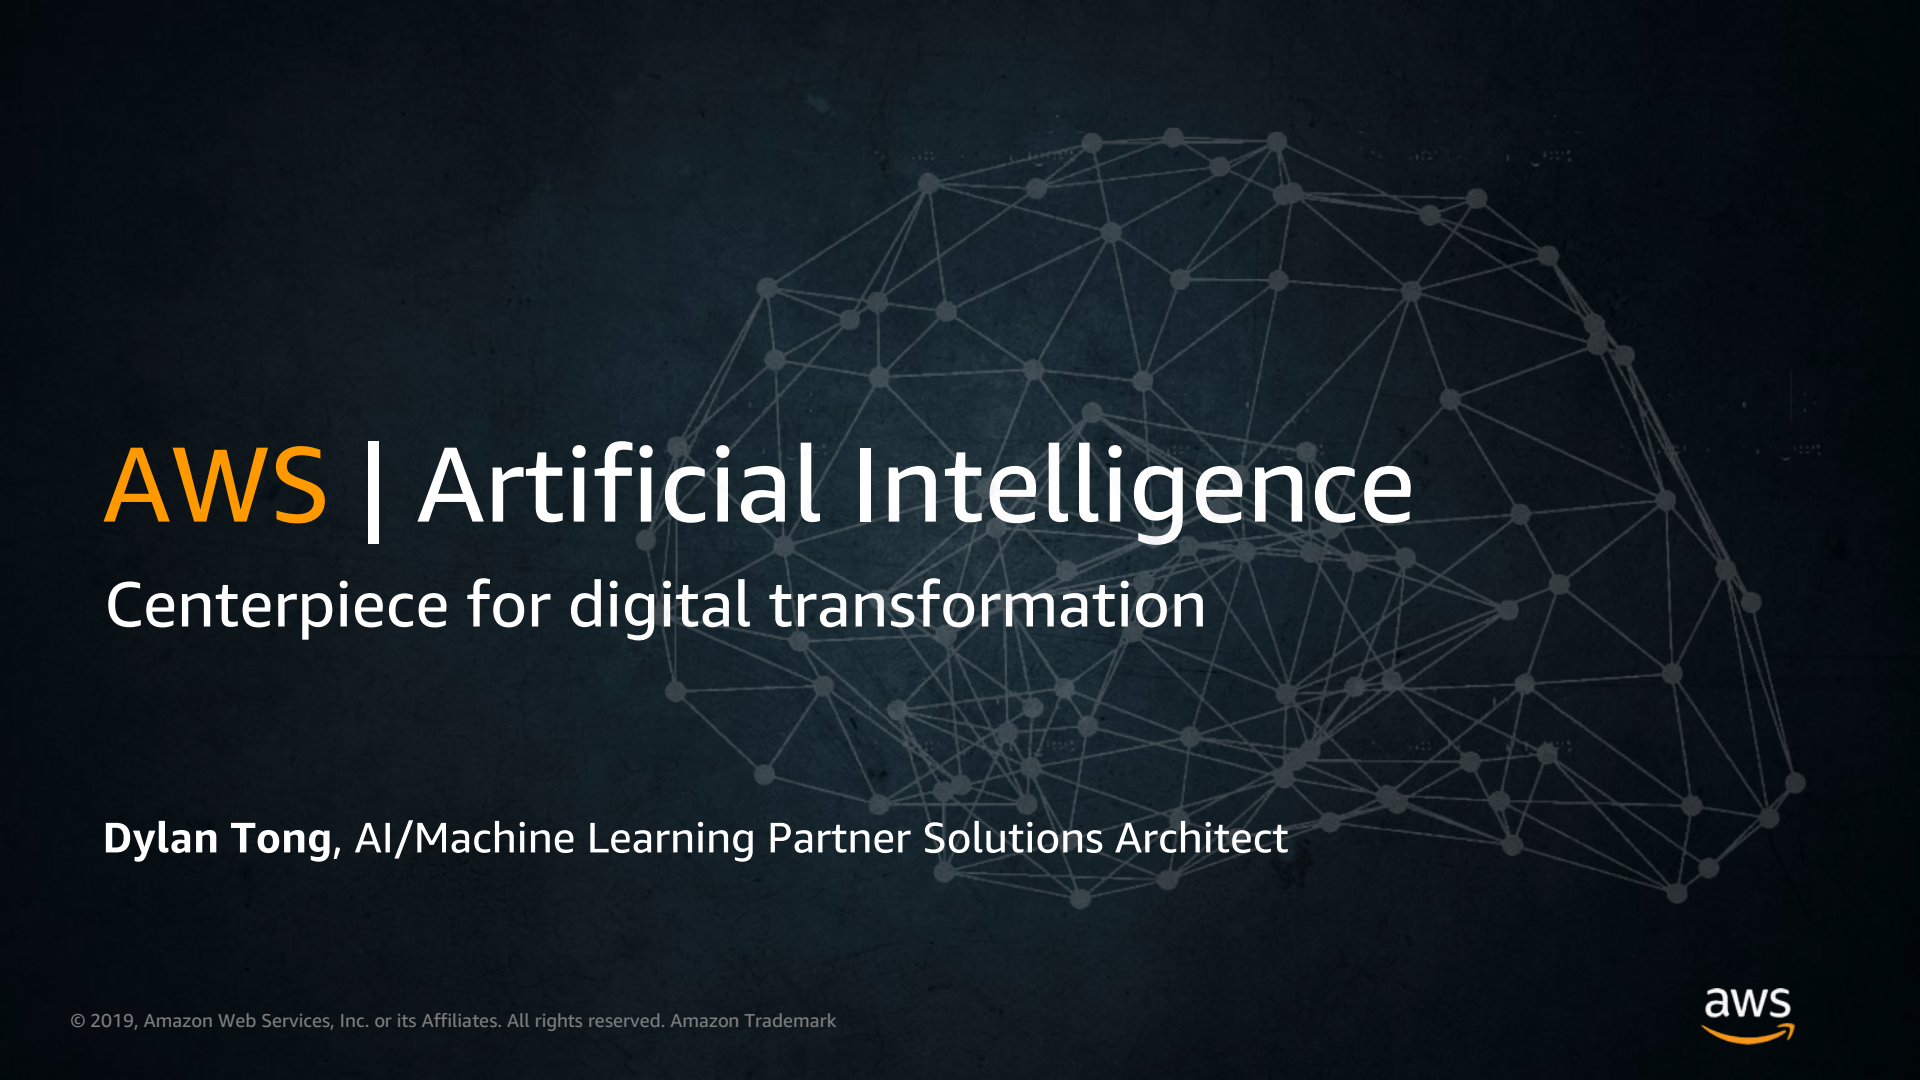

AWS | Artificial Intelligence

Centerpiece for digital transformation

Dylan Tong, AI/Machine Learning Partner Solutions Architect

What is AI?

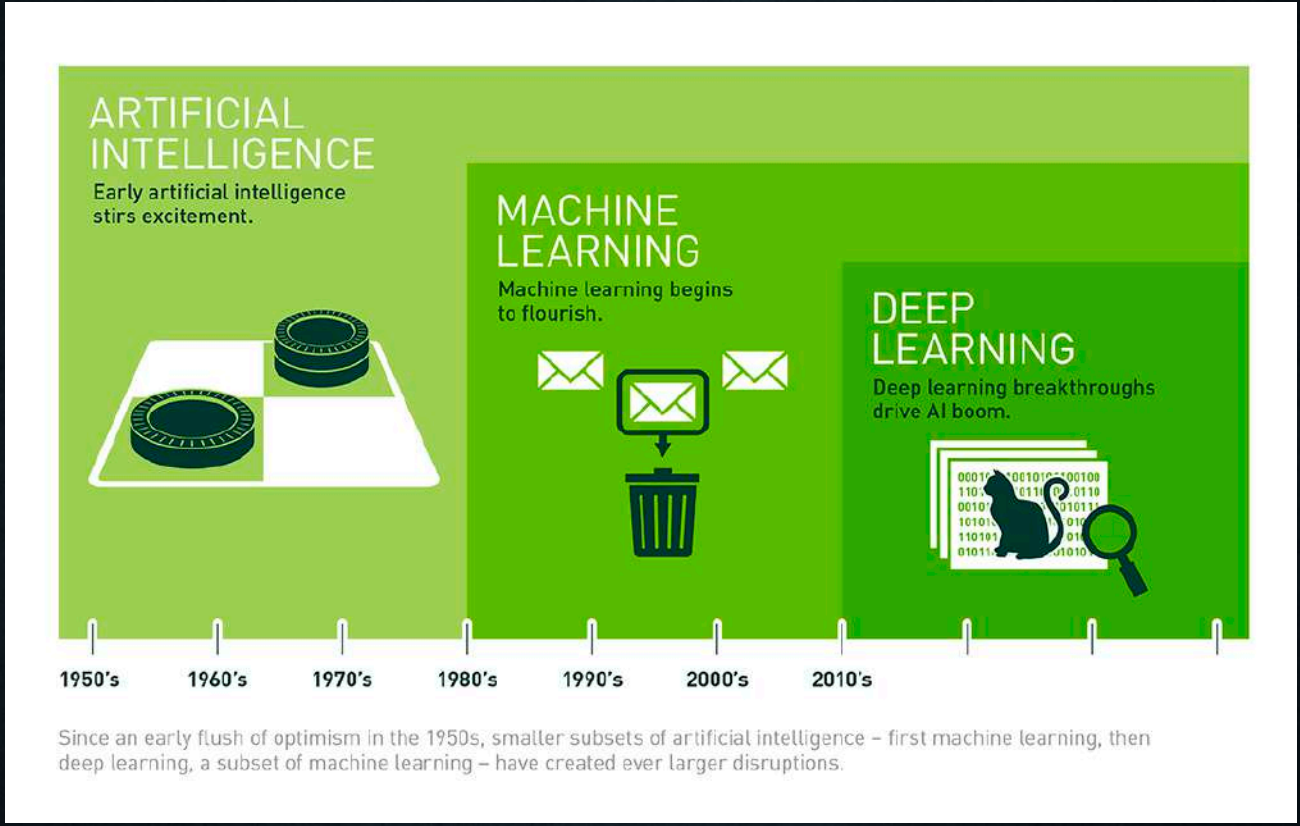

Reinventing the Retail Experience

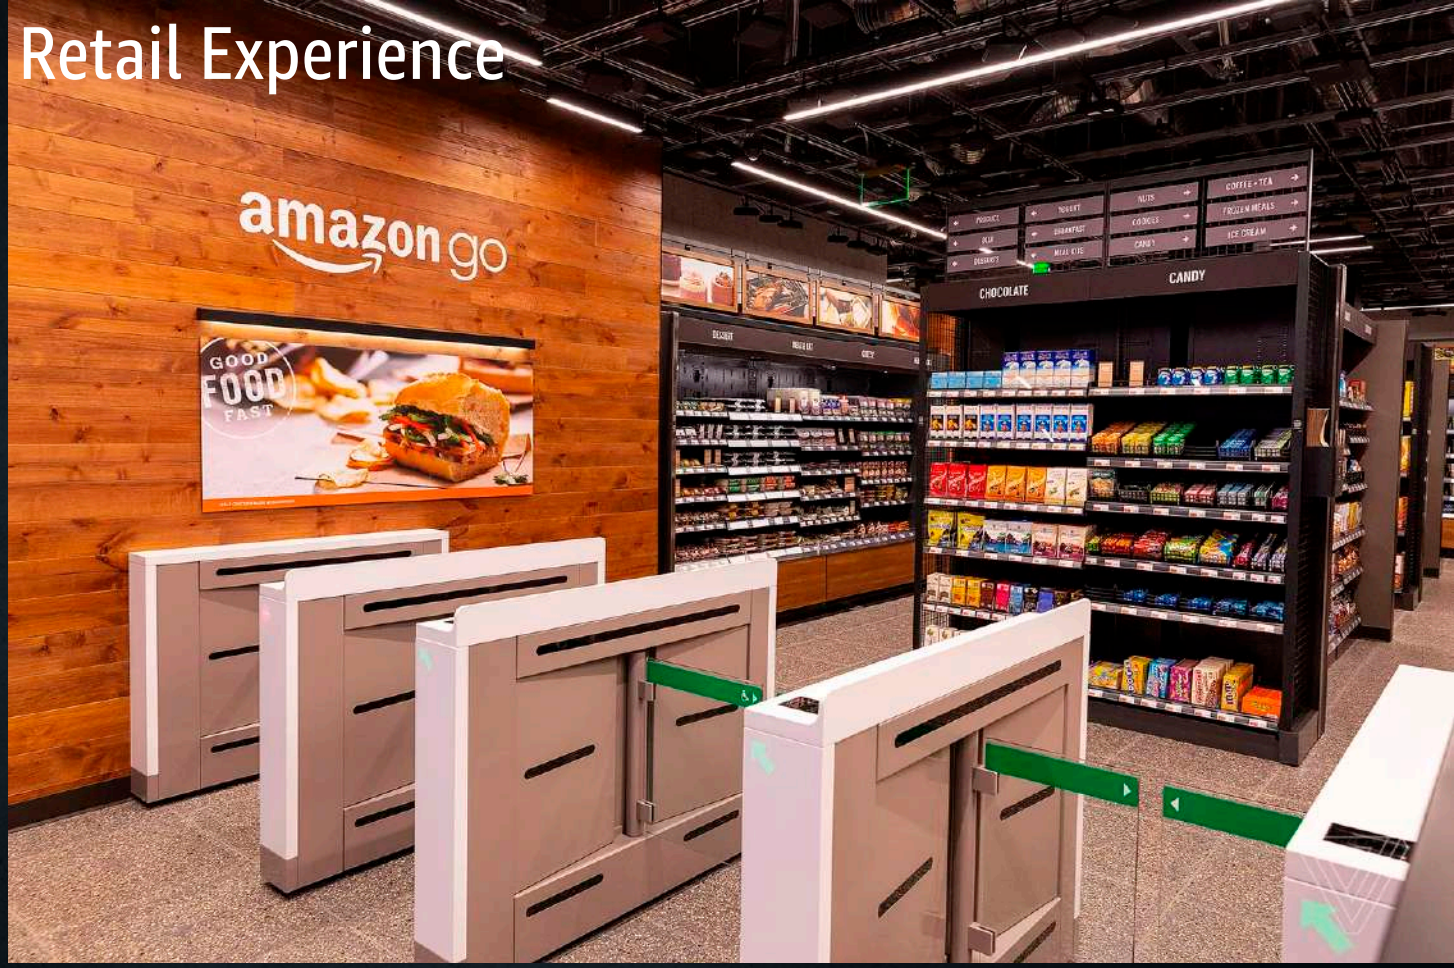

aws

AI Driven Stylist

INTRODUCING STYLE CHECK

Submit two photos to Style Check for a second opinion on which outfit looks better on you and why—based on fit, color, styling, and current trends. Through your feedback and input from our team of experienced fashion specialists, this advice gets smarter over time.

72% 28%

The styling of the pieces looks better.

aws

Smart Factory

Inspection:
Granular Quality and Progress Tracking

Asset Tracing (SKU)

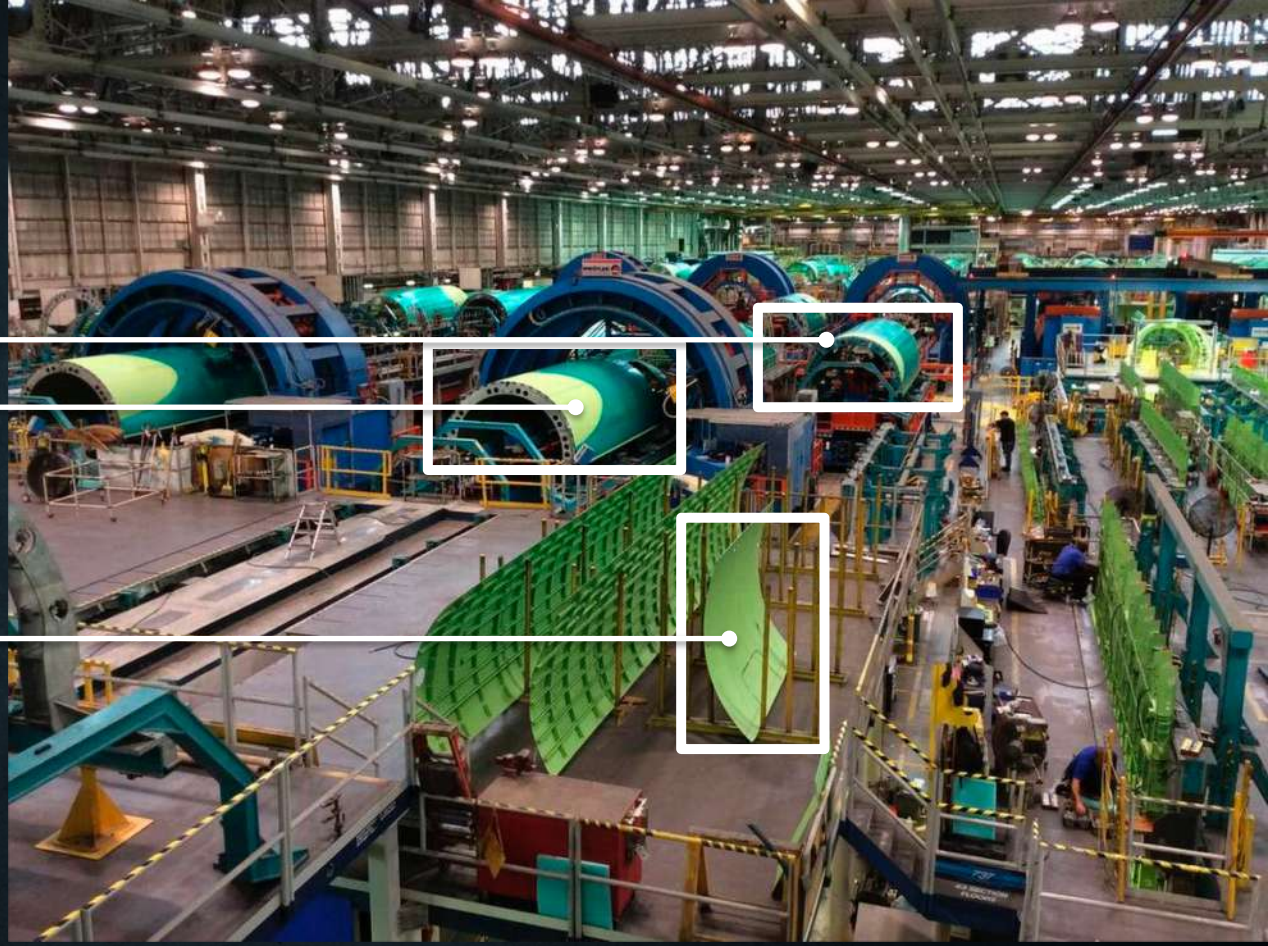

aws

The Connected Worker

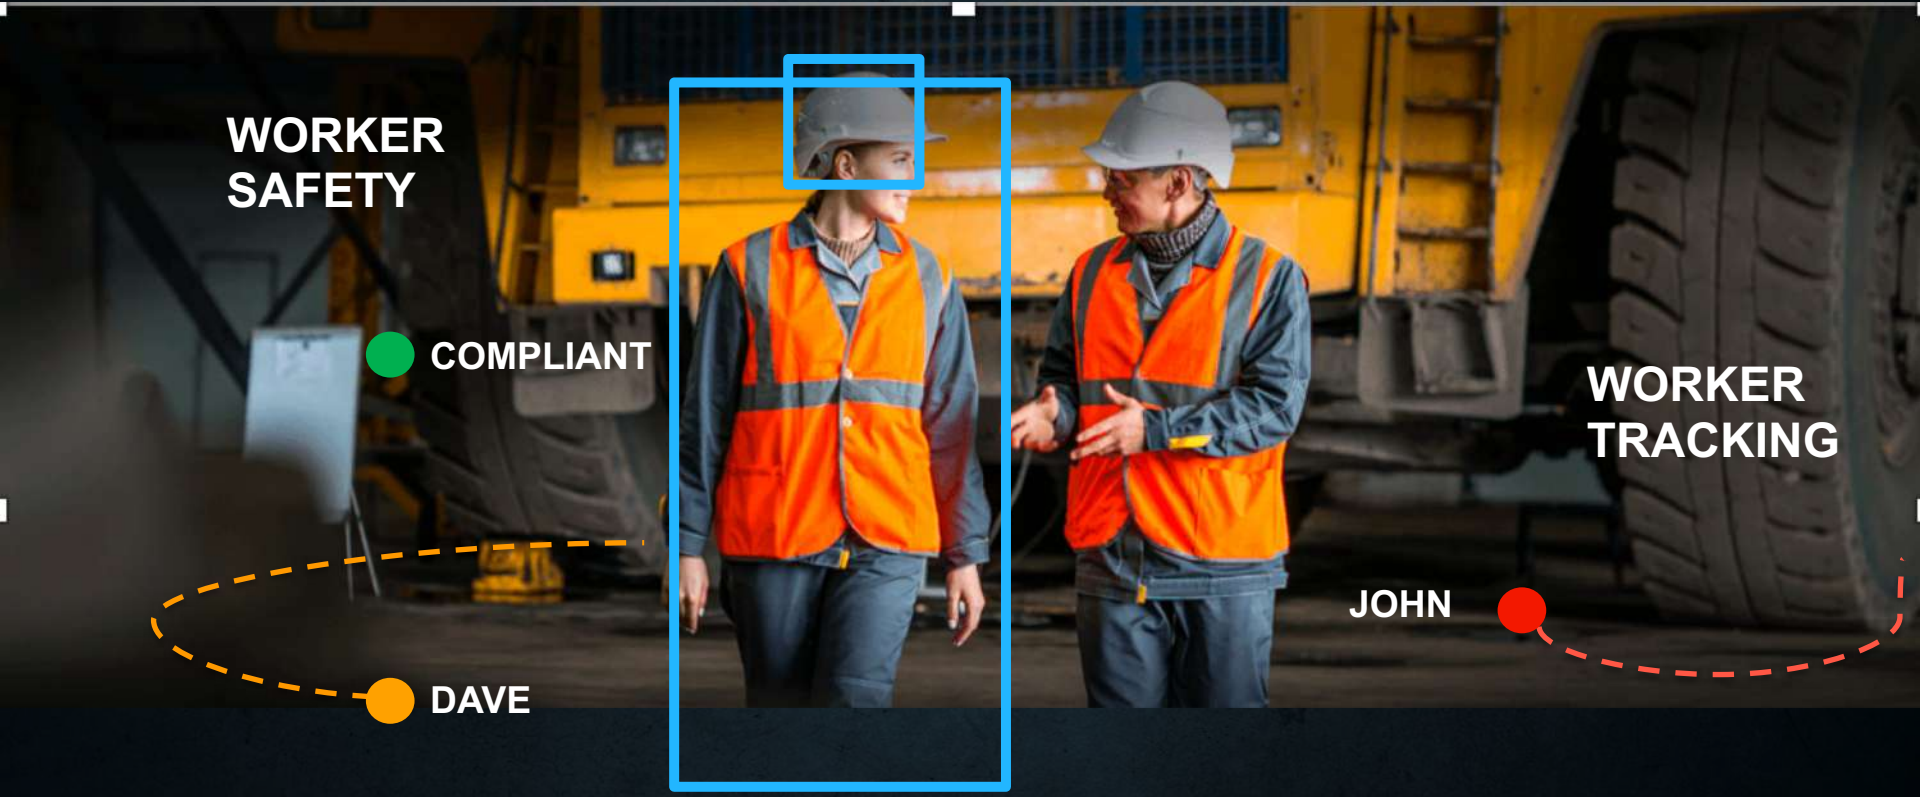

The Connected Worker

Improve worker productivity through new digital experiences:

- Voice-enabled Interfaces
- Augmented Reality

Reduce key-strokes and lower the learning curve of complex machinery.

JOHN DEERE

Smart Machine Era:

Reduces 90% use of herbicide through Blue River's smart sprayer. Computer Vision technology built on AWS and NVIDA GPU.

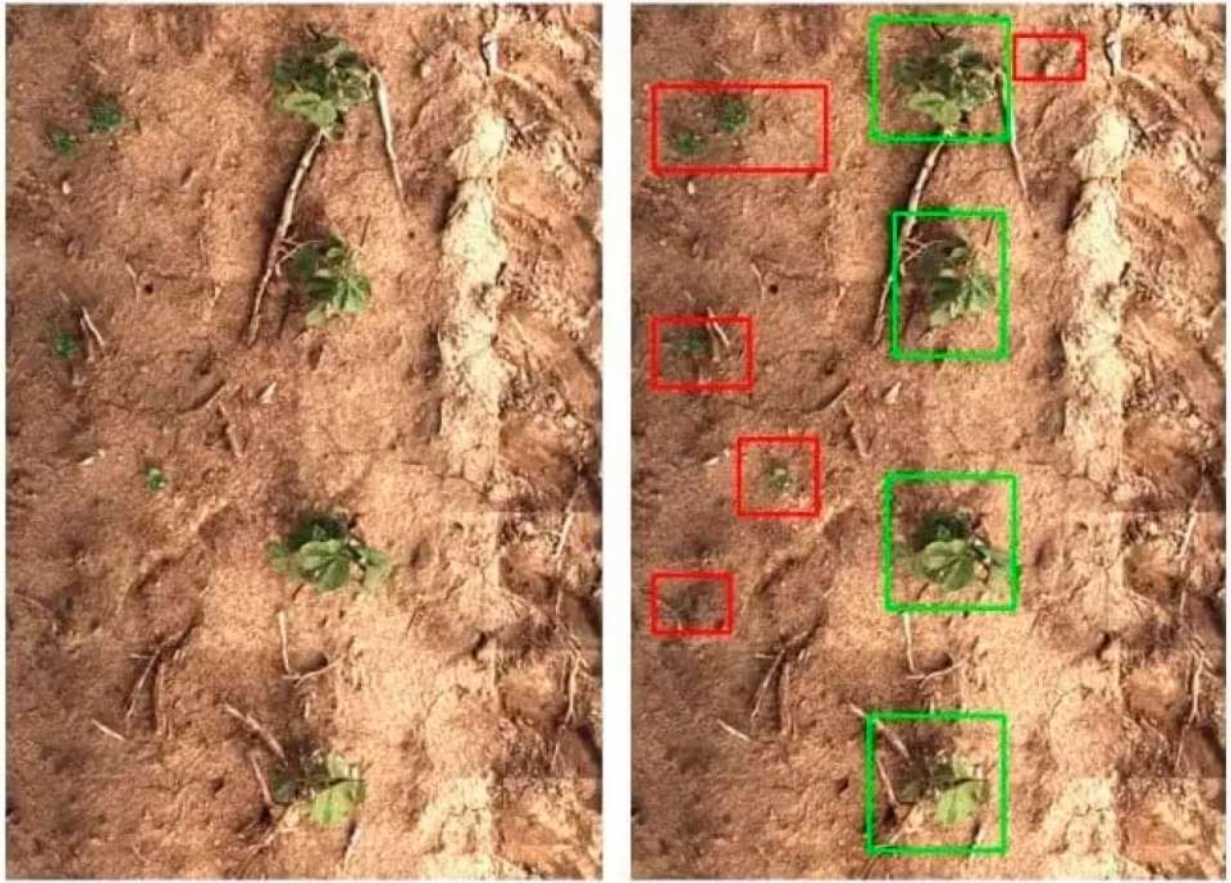

aws

DIUS

SOLVE GEOSOLUTIONS

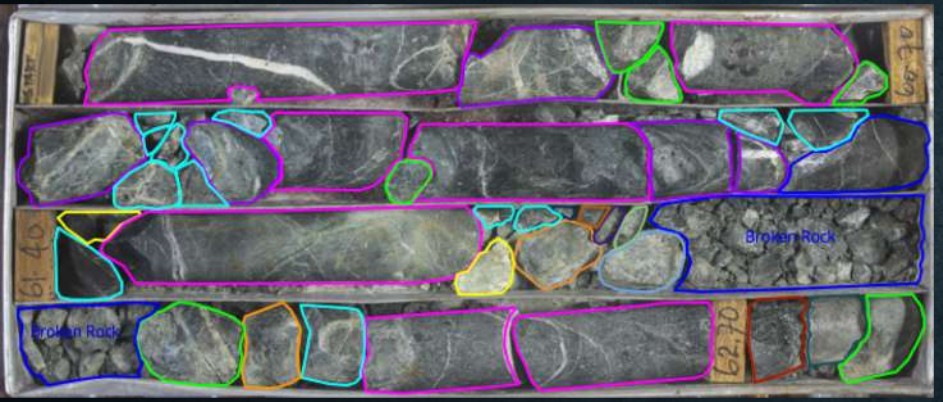

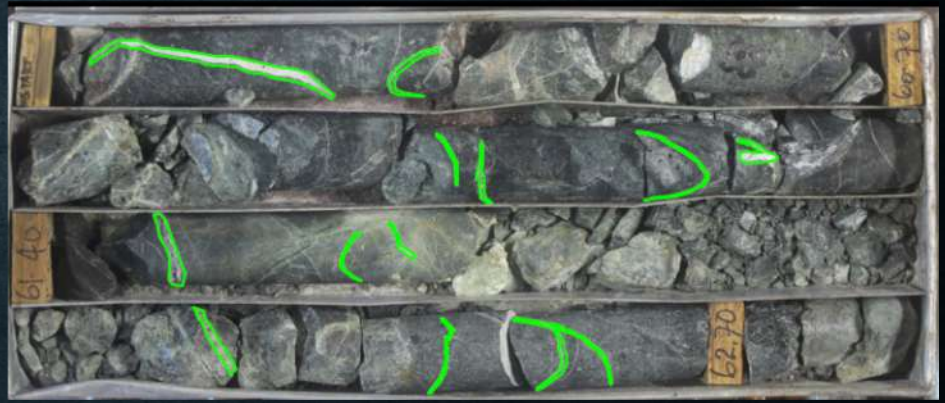

aws

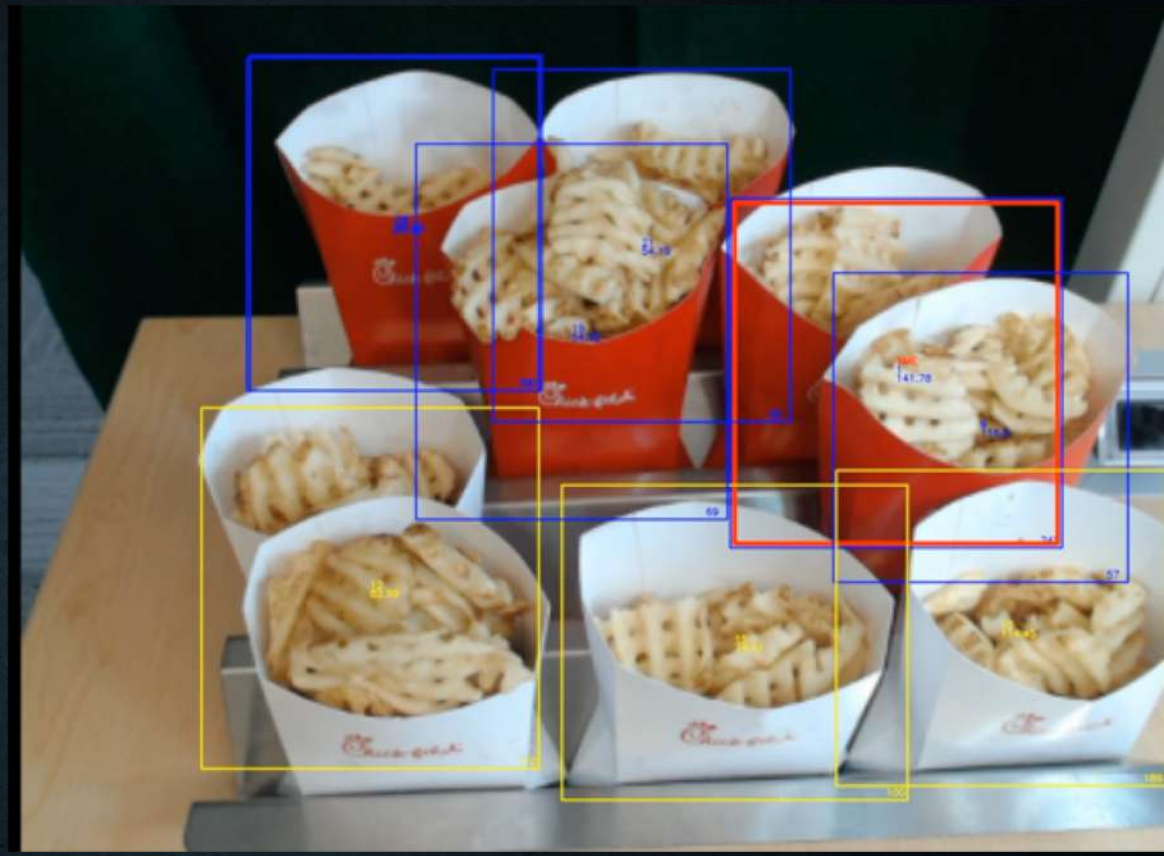

Technology

Mission: to put machine learning in the hands of every developer.

The Amazon ML stack: Broadest & deepest set of capabilities

AI SERVICES

Vision


REKOGNITION IMAGE


REKOGNITION VIDEO


TEXTRACT

Speech


POLLY


TRANSCRIBE

Language


TRANSLATE


COMPREHEND

Chatbots


LEX

Forecasting


FORECAST

Recommendations


PERSONALIZE

ML SERVICES


AMAZON SAGEMAKER

BUILD

Pre-built algorithms & notebooks

Data labeling (GROUND TRUTH)

Algorithms & models (AWS MARKETPLACE FOR MACHINE LEARNING)

TRAIN

One-click model training & tuning

Optimization (NEO)

Reinforcement learning

DEPLOY

One-click deployment & hosting

ML FRAMEWORKS & INFRASTRUCTURE

Frameworks

TensorFlow

mxnet

PYTORCH

Interfaces

GLUON

Keras

Infrastructure


EC2 P3 & P3N


EC2 C5


FPGAs


GREENGRASS


ELASTIC INFERENCE

aws

The Amazon ML stack: Frameworks and Infrastructure

aws

ML Optimized Environments and Infrastructure

More data science, less setup

Deep Learning AMI

aws

P3dn Instances: Optimized for Training at Scale

Reduce machine learning training time

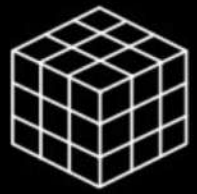

Better GPU utilization

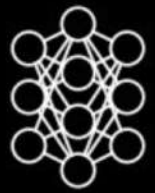

Support larger, more complex models

KEY FEATURES

100Gbps of networking bandwidth (4x more P3)

8 NVIDIA Tesla V100 GPUs

32GB of memory per GPU (total 256GB, 2x more P3)

96 Intel Skylake vCPUs (50% more than P3) with AVX-512

aws

Challenges with Inference in Production

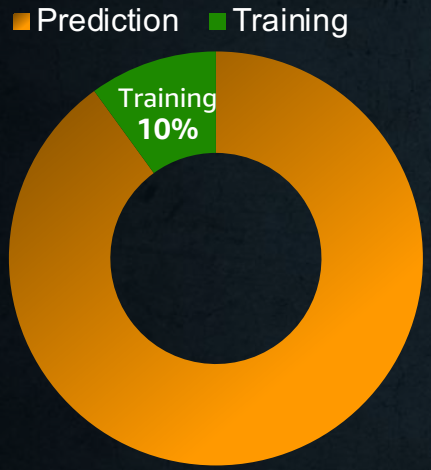

Low utilization and high costs

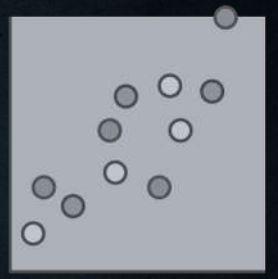

One size does not fit all

aws

Amazon Elastic Inference

Reduce Deep Learning Inference costs up to 75%

Lower inference costs

Match capacity
to demand

Available between 1 to 32 TFLOPS
per accelerator

KEY FEATURES

Integrated with
Amazon EC2 and
Amazon SageMaker

Support for TensorFlow, Apache
MXNet, and ONNX
with PyTorch coming soon

Single and
mixed-precision
operations

aws

Up to 75% Reduction in Inference Costs

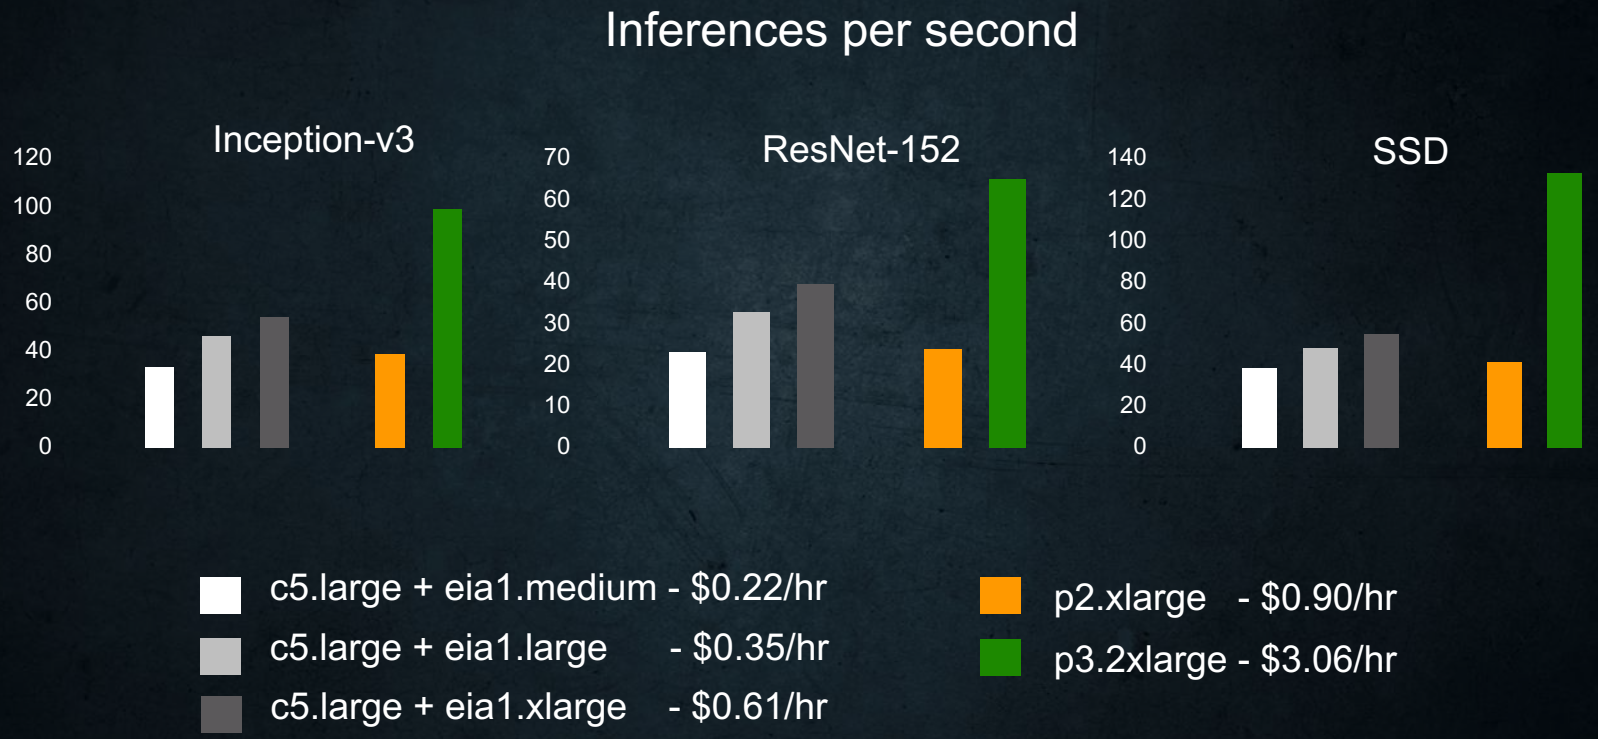

Requirements for Inference at the Edge

BANDWIDTH

1 billion cameras WW (2020)
10's of petabytes per day

LATENCY

30 images per second
200ms latency

PRIVACY

Confidentiality
Private cloud or on-premises storage

AVAILABILITY

50% of populated world < 8mbps
Bulk of uninhabited world no 3G+

aws

AWS Greengrass

Extend intelligence to the edge

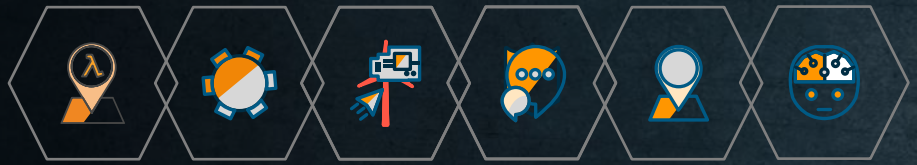

Local actions | Local triggers | Data and state sync | Security | Local resource access | Local ML inference

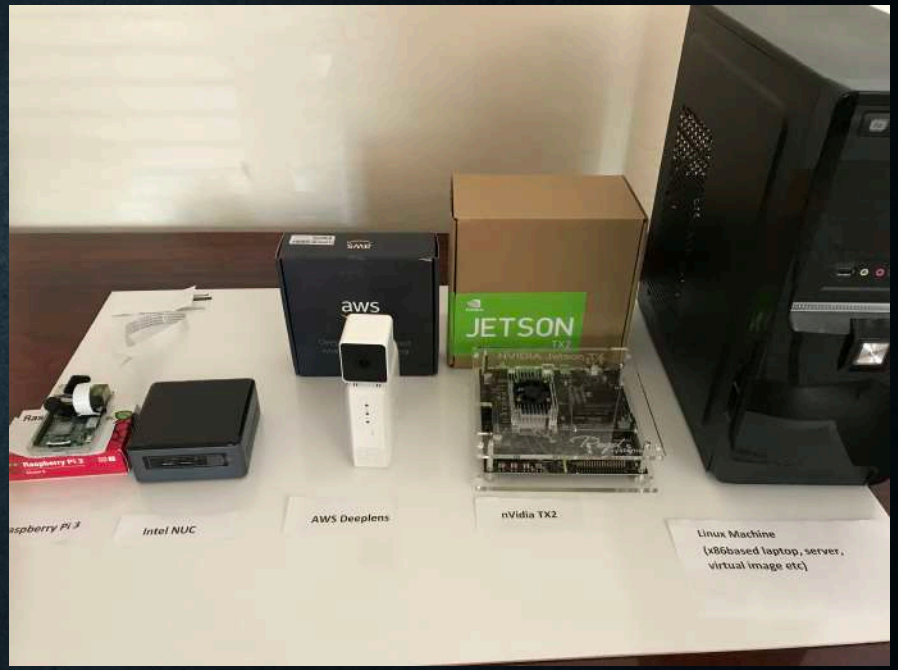

aws

The Amazon ML stack: ML Services

ML SERVICES

AMAZON SAGEMAKER

BUILD	TRAIN	DEPLOY
Pre-built algorithms & notebooks	One-click model training & tuning	One-click deployment & hosting
Data labeling (GROUND TRUTH)	Optimization (NEO)	
Algorithms & models (AWS MARKETPLACE FOR MACHINE LEARNING)	Reinforcement learning	

ML FRAMEWORKS & INFRASTRUCTURE

Frameworks

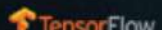

Interfaces

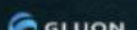

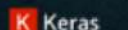

Infrastructure

EC2 P3 & P3N

EC2 C5

FPGAs

GREENGRASS

ELASTIC INFERENCE

aws

Amazon SageMaker

BRINGING MACHINE LEARNING TO ALL DEVELOPERS

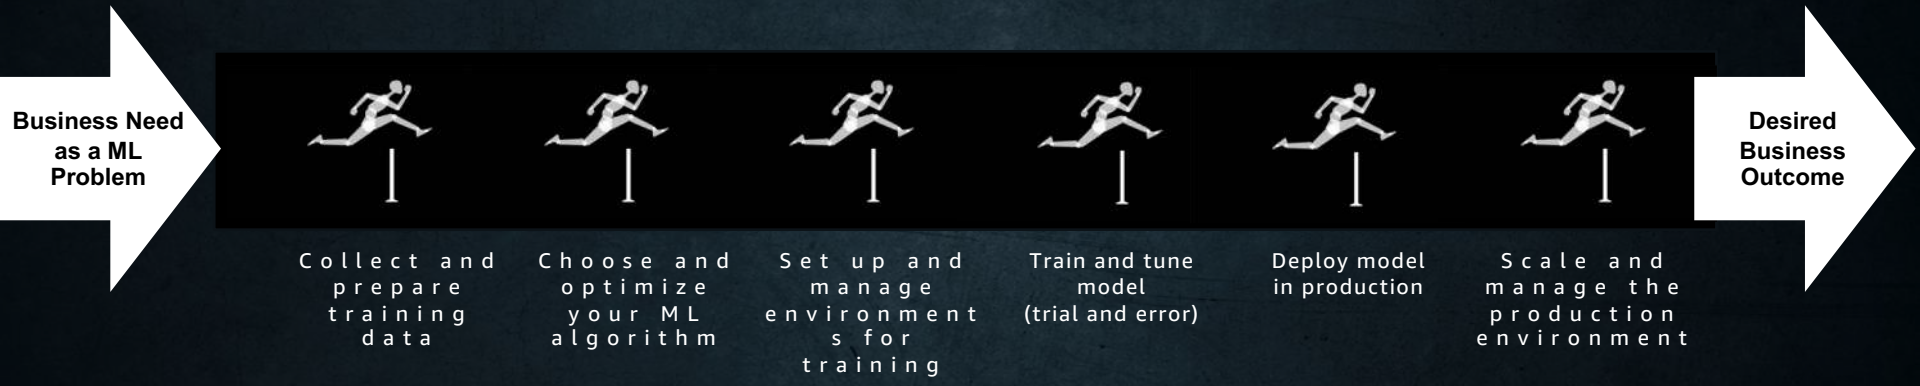

SIMPLIFY THE END-TO-END MACHINE LEARNING PROCESS

aws

Training Data for Supervised Learning

Challenges

- Timely annotation of large data sets
- Managing workforces.
  - Ensuring labeling quality
  - Handling workflows
- Integration with ML development environment

aws

Amazon SageMaker Ground Truth

Label machine learning training data easily and accurately

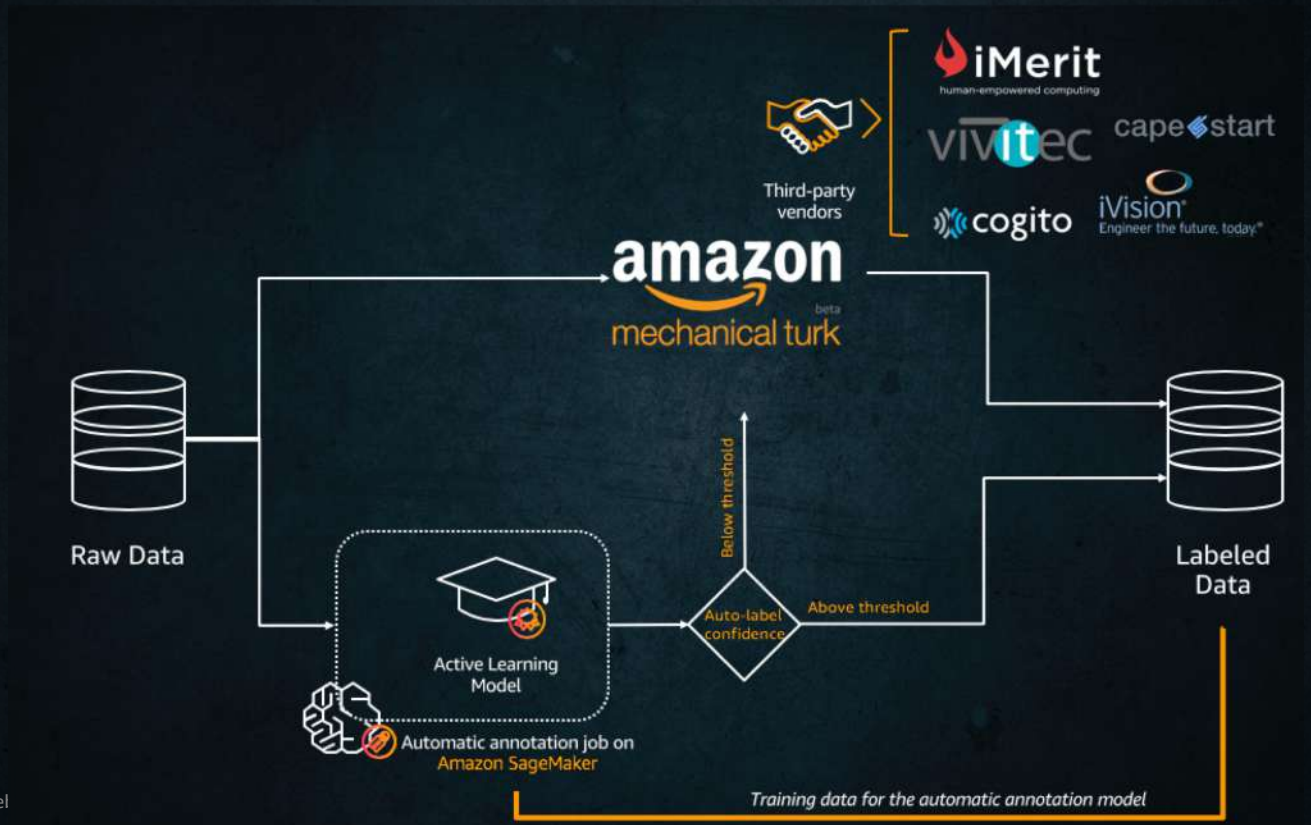

aws

Managed Notebooks: Exploration, Experimentation

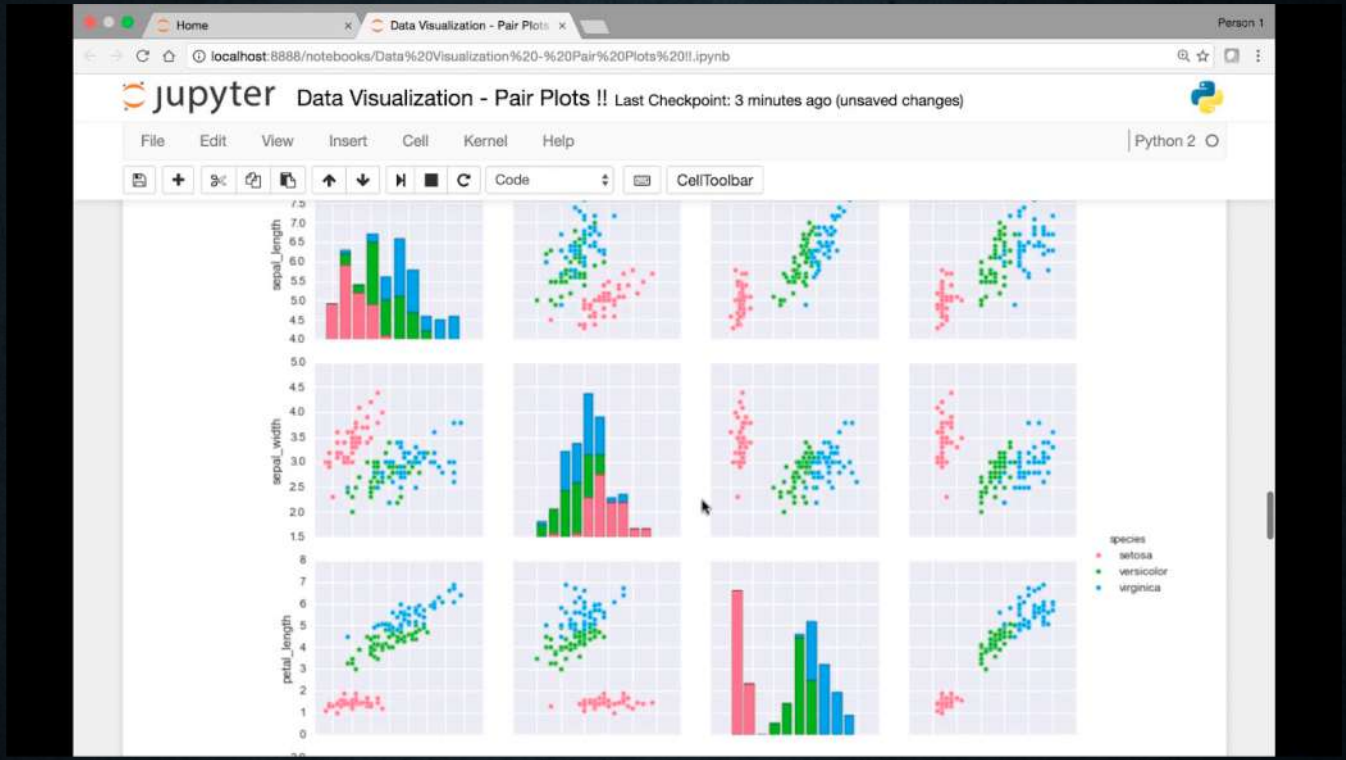

aws

Your Choice. The Right Tools.

NATIVE SUPPORT FOR MOST POPULAR FRAMEWORKS

SAGEMAKER OPTIMIZED ALGORITHMS

- BlazingText Algorithm
- DeepAR Forecasting Algorithm
- Factorization Machines Algorithm
- Image Classification Algorithm
- IP Insights Algorithm
- K-Means Algorithm
- K-Nearest Neighbors (k-NN) Algorithm
- Latent Dirichlet Allocation (LDA) Algorithm
- Linear Learner Algorithm
- Neural Topic Model (NTM) Algorithm
- Object2Vec Algorithm
- Object Detection Algorithm
- Principal Component Analysis (PCA) Algorithm
- Random Cut Forest (RCF) Algorithm
- Semantic Segmentation Algorithm
- Sequence-to-Sequence Algorithm
- XGBoost Algorithm

aws marketplace INTEGRATION

DEEP VISION
H2O.ai
figure eight
TIBCO
D.M DIMENSIONAL MECHANICS
Open for Innovation KNIME
wipro
Knowledgent

aws

Zero Setup Training

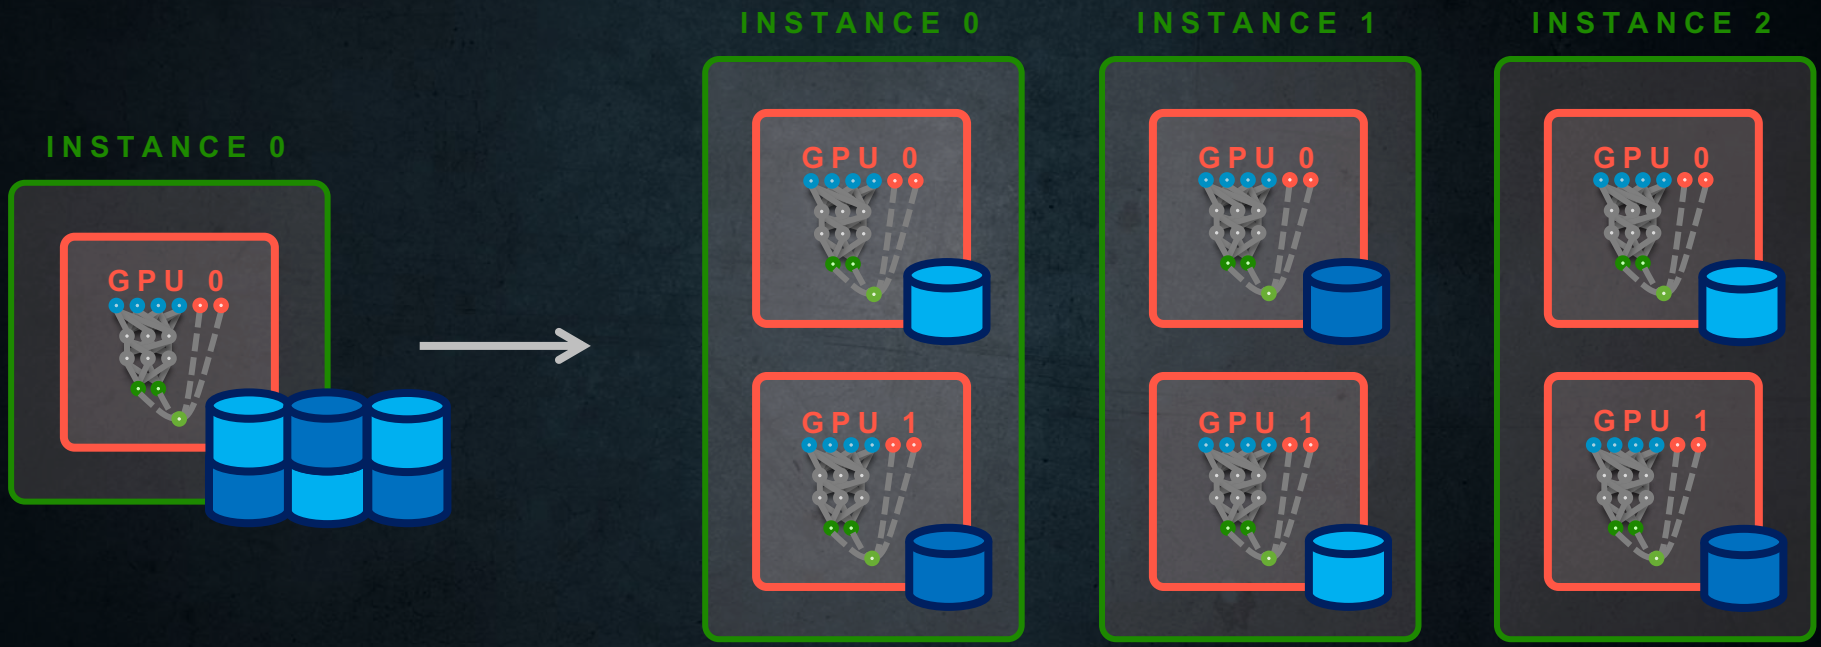

aws

Automatic Model Tuning

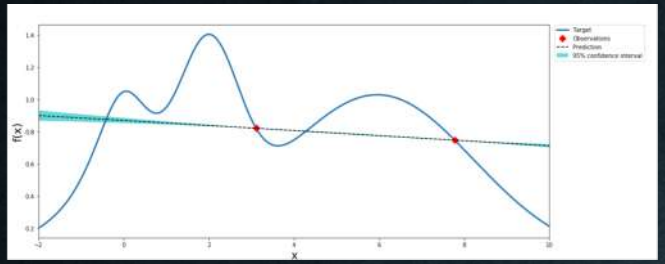

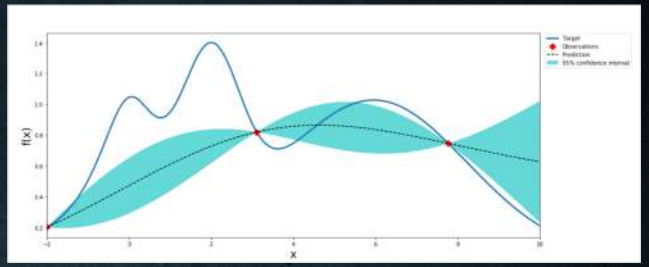

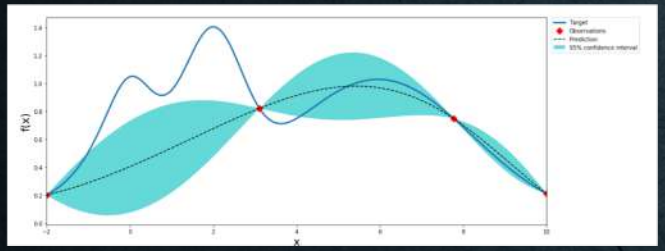

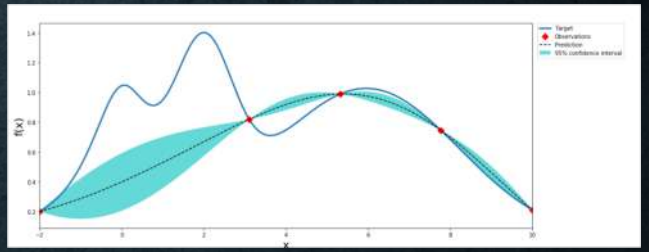

. . .

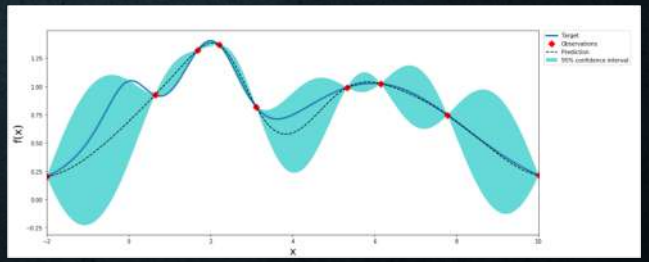

aws

Amazon SageMaker Neo

Train once, run everywhere with 2x performance and no accuracy lost

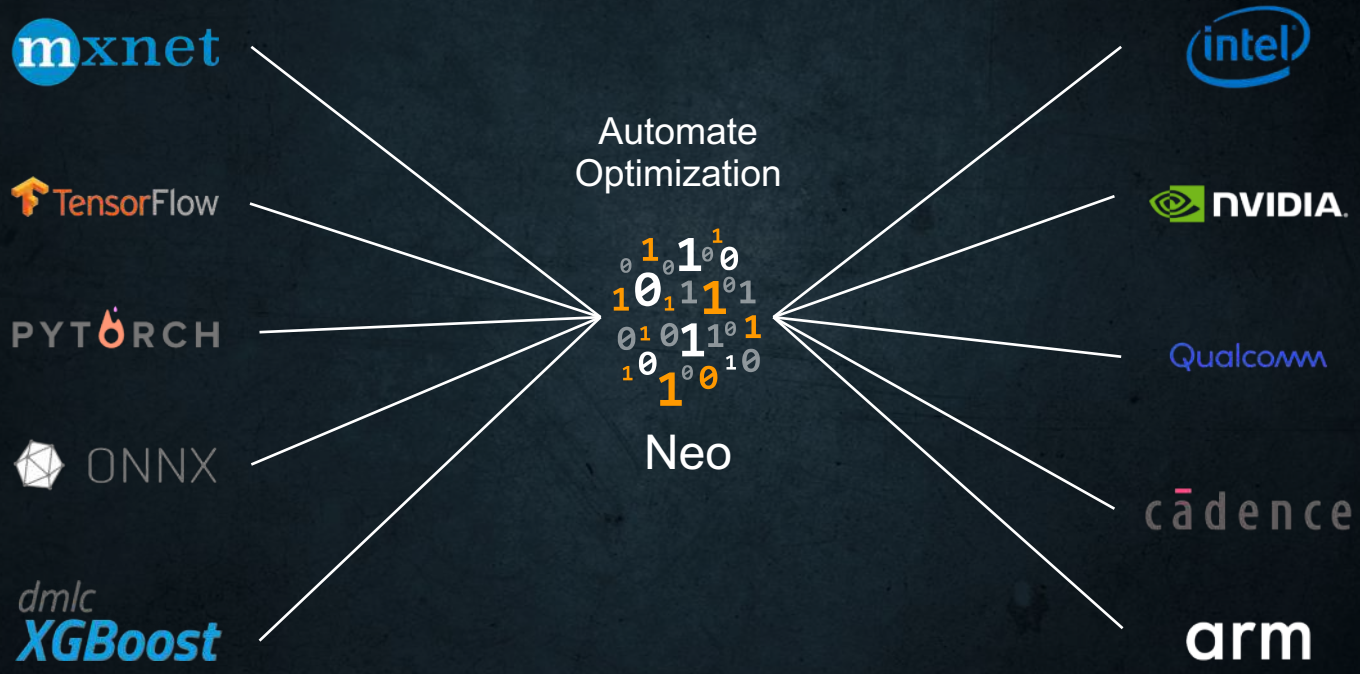

aws

Simplified Model Deployment and Hosting

SageMaker Hosted Endpoints

- Auto-scaling
- Performance monitoring
- A/B Testing
- Elastic Inference support
- Suited for real-time and batch workloads

aws

Custom machine learning for your business

AMAZON SAGEMAKER

REDUCE COSTS

70%

cost reduction for data labeling using Ground Truth

75%

cost reduction for inference with Elastic Inference

INCREASE PERFORMANCE

10x

better algorithm performance

2x

performance increases from model optimization with Neo

EASE-OF-USE

One-click

model training & deployment

Train once

run anywhere

aws

The Amazon ML stack: Broadest & deepest set of capabilities

AI SERVICES

Vision


REKOGNITION IMAGE


REKOGNITION VIDEO


TEXTRACT

Speech


POLLY


TRANSCRIBE

Language


TRANSLATE


COMPREHEND

Chatbots


LEX

Forecasting


FORECAST

Recommendations


PERSONALIZE

ML SERVICES


AMAZON SAGEMAKER

BUILD

Pre-built algorithms & notebooks

Data labeling (GROUND TRUTH)

Algorithms & models (AWS MARKETPLACE FOR MACHINE LEARNING)

TRAIN

One-click model training & tuning

Optimization (NEO)

Reinforcement learning

DEPLOY

One-click deployment & hosting

ML FRAMEWORKS & INFRASTRUCTURE

Frameworks

TensorFlow

mxnet

PYTORCH

Interfaces

GLUON

Keras

Infrastructure


EC2 P3 & P3N


EC2 C5


FPGAs


GREENGRASS


ELASTIC INFERENCE

aws

Chatbots with Amazon Lex

Speech to Intent

ASR+NLU integrated into one API

Text to Speech

Amazon Polly integrated into API

Dialog Management

Session context maintenance

Deployment

One click deployment

Scale

Completely managed service

Business Logic

Native integration with AWS Lambda

Security

Encrypted data in transit & at rest

Analytics

Monitor and improve

End to End

Vision: Amazon Rekognition

IMAGES

Object and Scene Detection

Facial Analysis

Face Recognition

Unsafe Image Detection

Celebrity Recognition

Text in Image

VIDEO

Person Tracking

Real-time Live Stream

aws

Vision: Amazon Rekognition

IMAGES

Object and Scene Detection

Facial Analysis

Face Recognition

Unsafe Image Detection

Celebrity Recognition

Text in Image

VIDEO

Person Tracking

Real-time Live Stream

aws

The Connected Worker

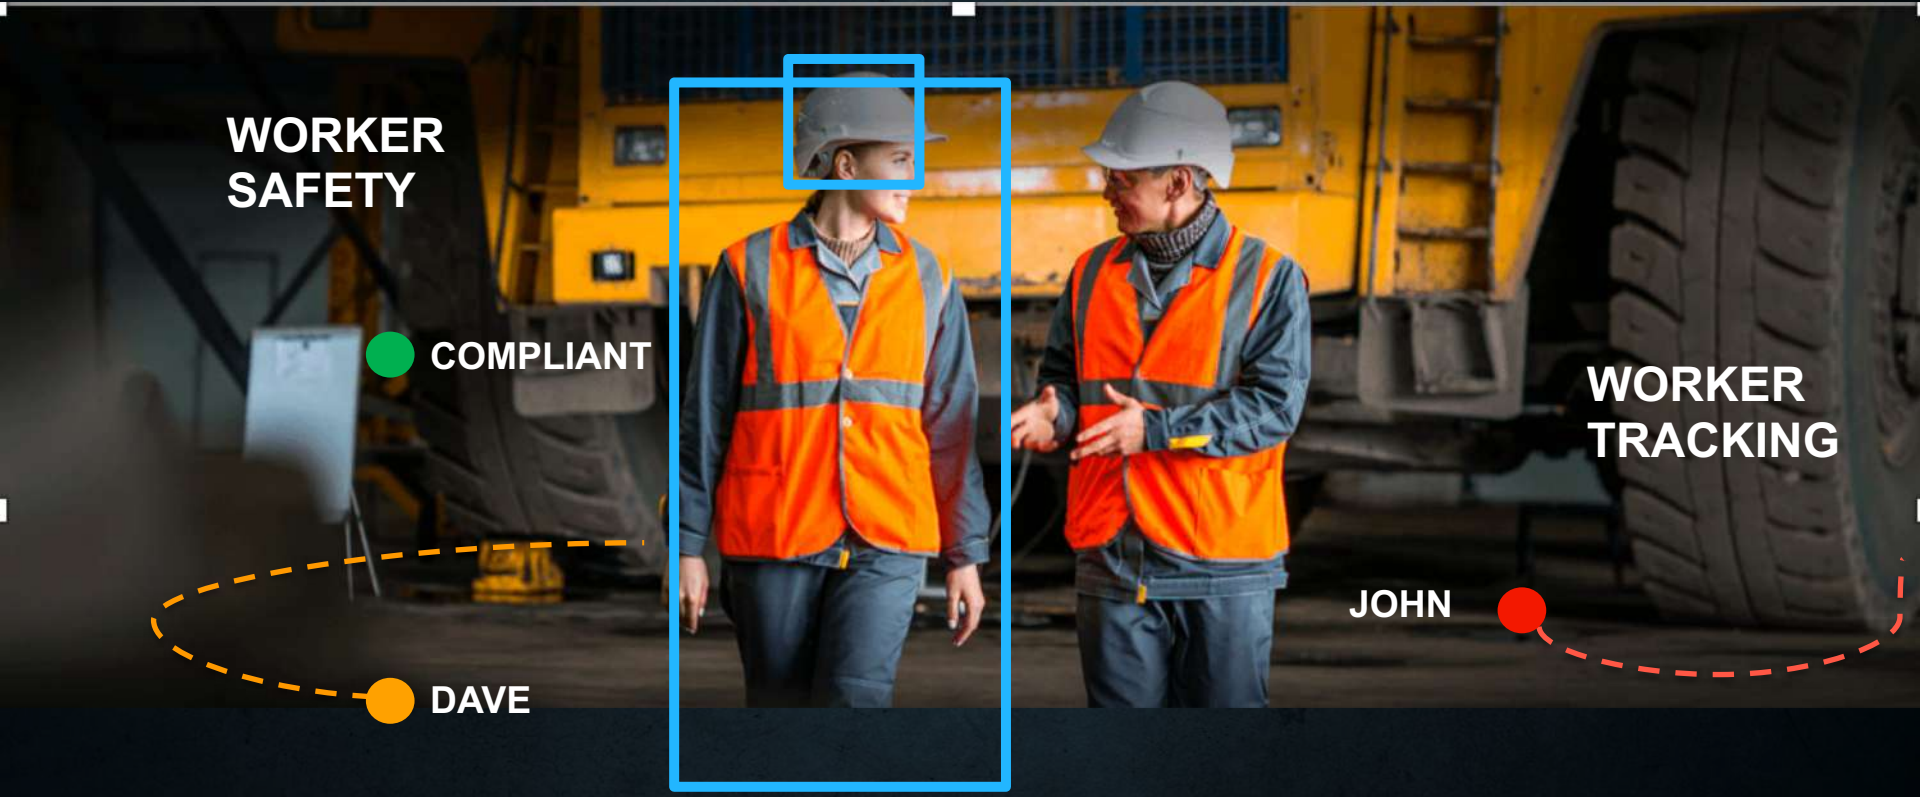

Vision: Amazon Textract

Form Extraction simplified

Full Name			Date of Birth			Gender	
John	X	Doe	01	01	1971	Male	●
First	Middle	Last	MM	DD	YYYY	Female	○

- ✓ Logical groupings captured
- ✓ Relationships captured
- ✓ Glyphs captured

Output

Full Name:
- First: John
- Middle: X
- Last: Doe

Date of Birth:
- MM: 01
- DD: 01
- YYYY: 1971

Gender:
- Male: True
- Female: False

aws

NLP: Amazon Comprehend

Amazon.com, Inc. is located in Seattle, WA and was founded July 5th, 1994 by Jeff Bezos. Our customers *love buying everything* from books to blenders at great prices

Named Entities

- Amazon.com: Organization
- Seattle, WA : Location
- July 5th, 1994: Date
- Jeff Bezos : Person

Keyphrases

- Our customers
- books
- blenders
- great prices

Sentiment

- *Positive*

Language

- English

aws

More Natural Language Processing…

TEXT-TO-SPEECH

Amazon Polly

SPEECH-TO-TEXT

Amazon Transcribe

LANGUAGE TRANSLATION

Amazon Translate

aws

Solution Acceleration: Intelligent Building Blocks

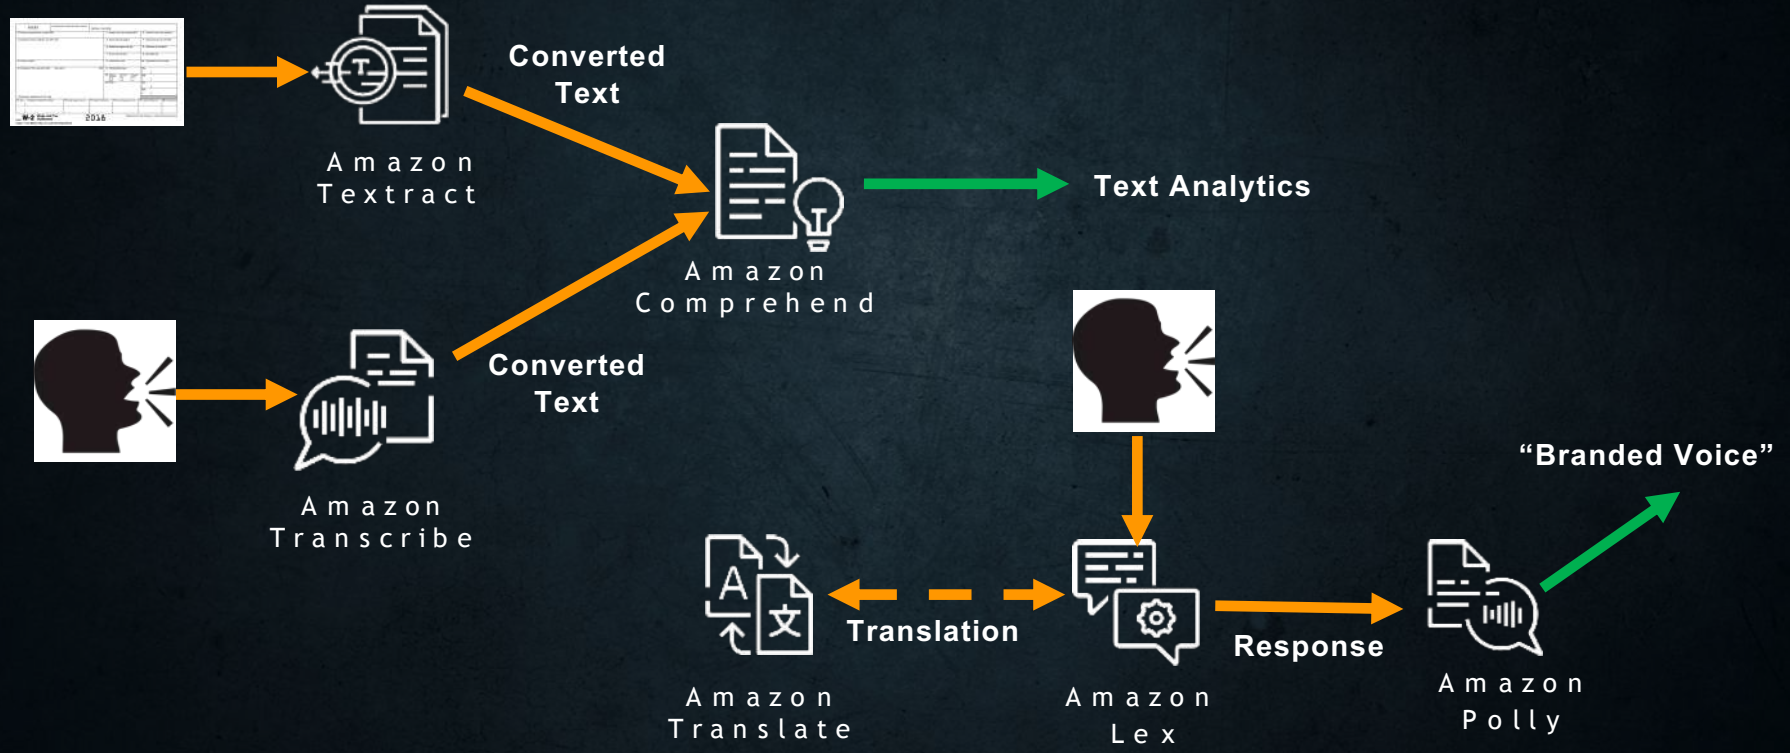

AutoML: Tailored Models Automated

FORECASTING

Amazon Forecast

RECOMMENDATIONS

Amazon Personalize

aws

Computer Vision on AWS Jumpstart

Jumpstart Workshop Menu

- Object Detection
- TBD…

Object Detection Workshop

Reference Architectures

Cloud Inference on Streaming Video

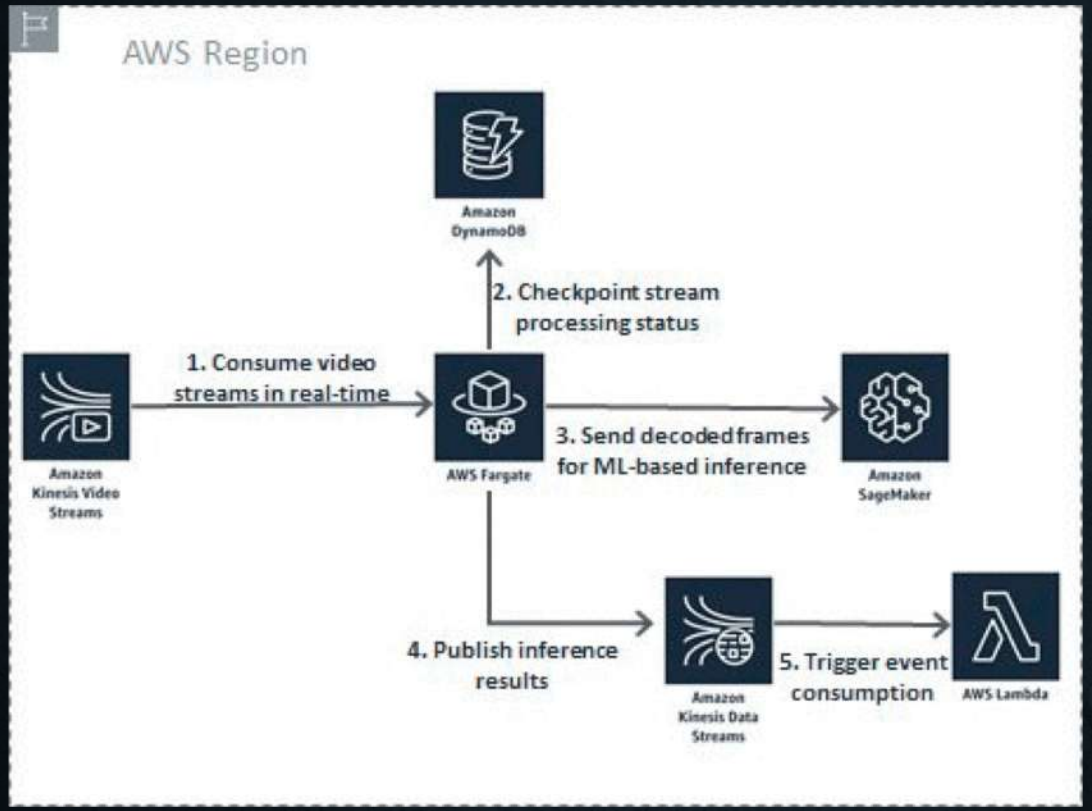

AWS industrial IoT reference architecture

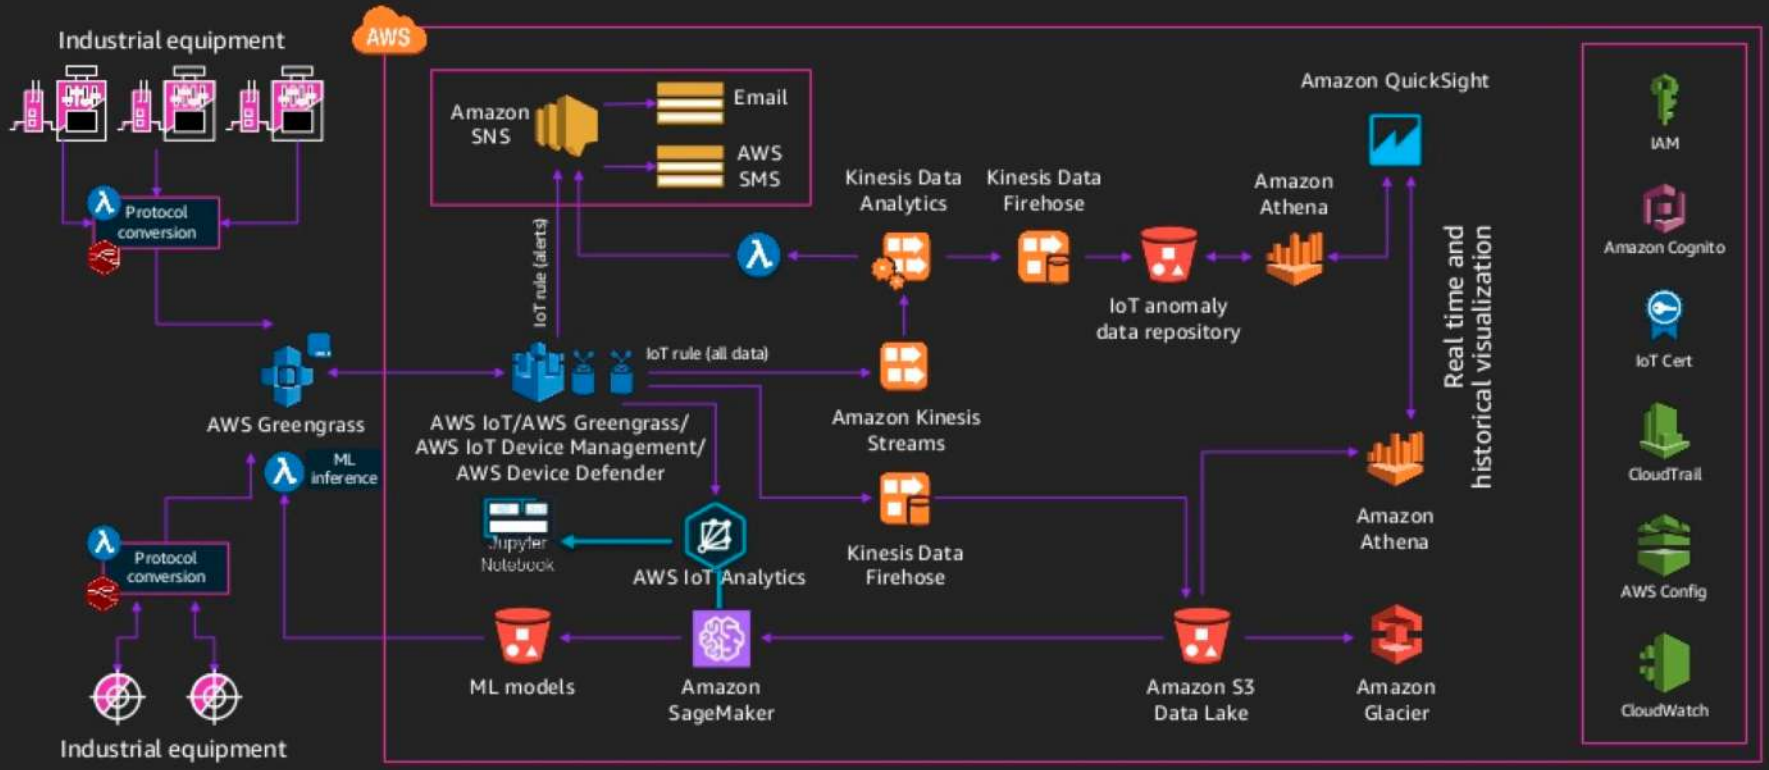

Workshop Objectives and Agenda

Select the Right Strategy

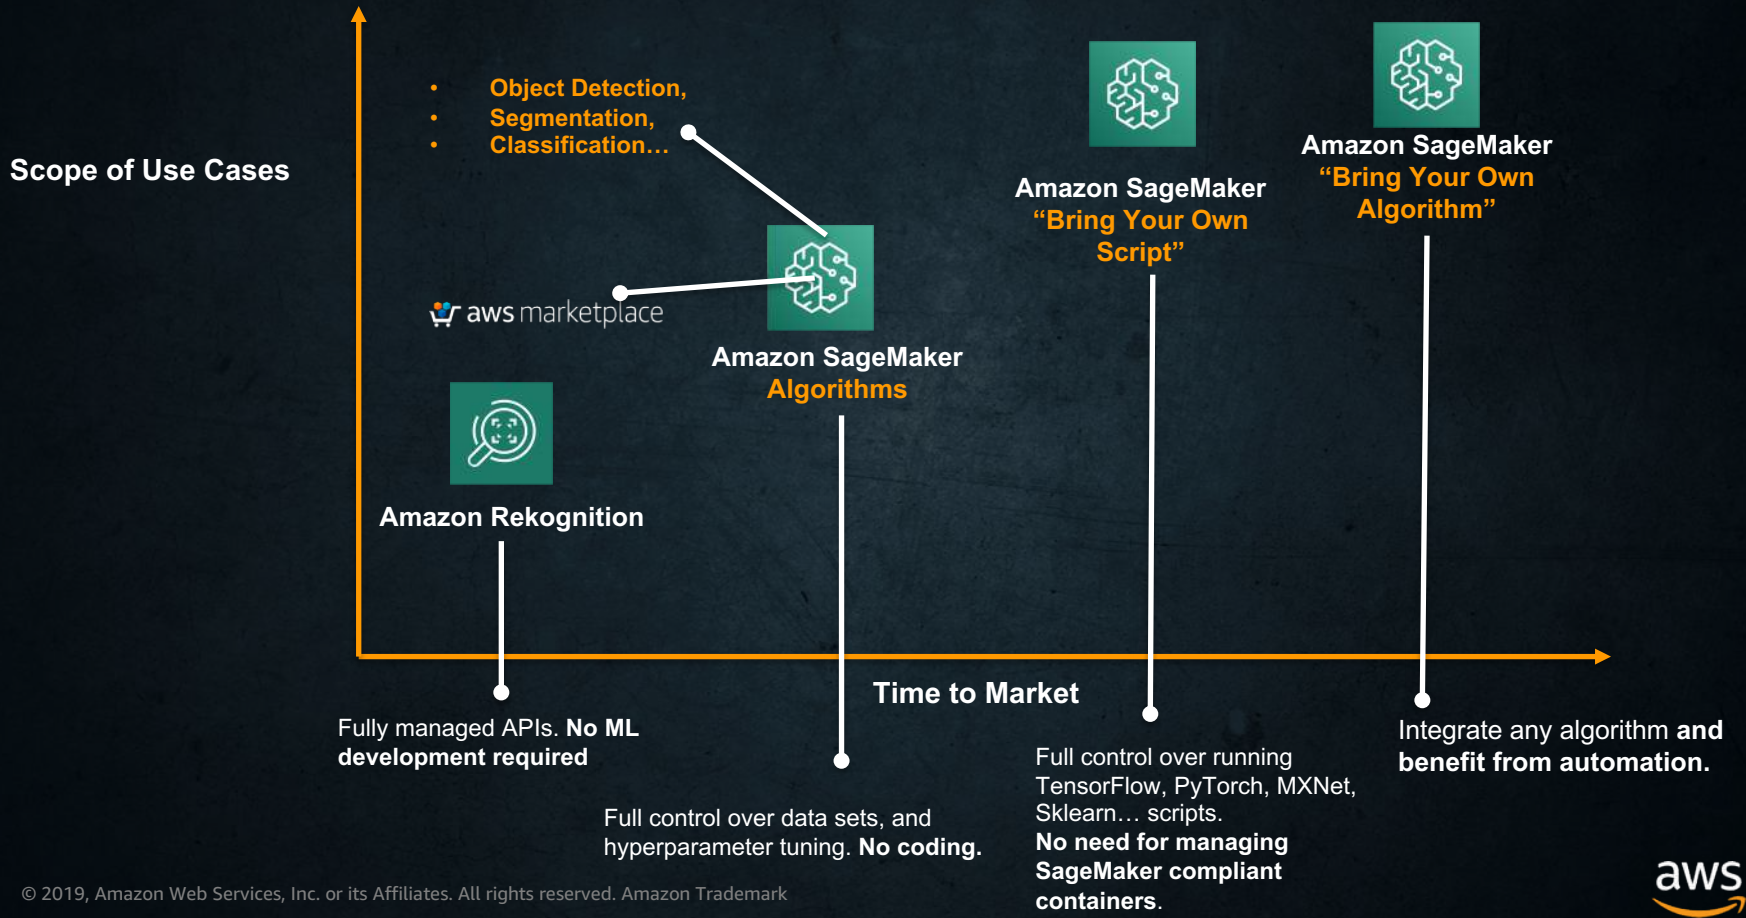

Right Tool for the Use Case

aws

Amazon SageMaker

BRINGING MACHINE LEARNING TO ALL DEVELOPERS

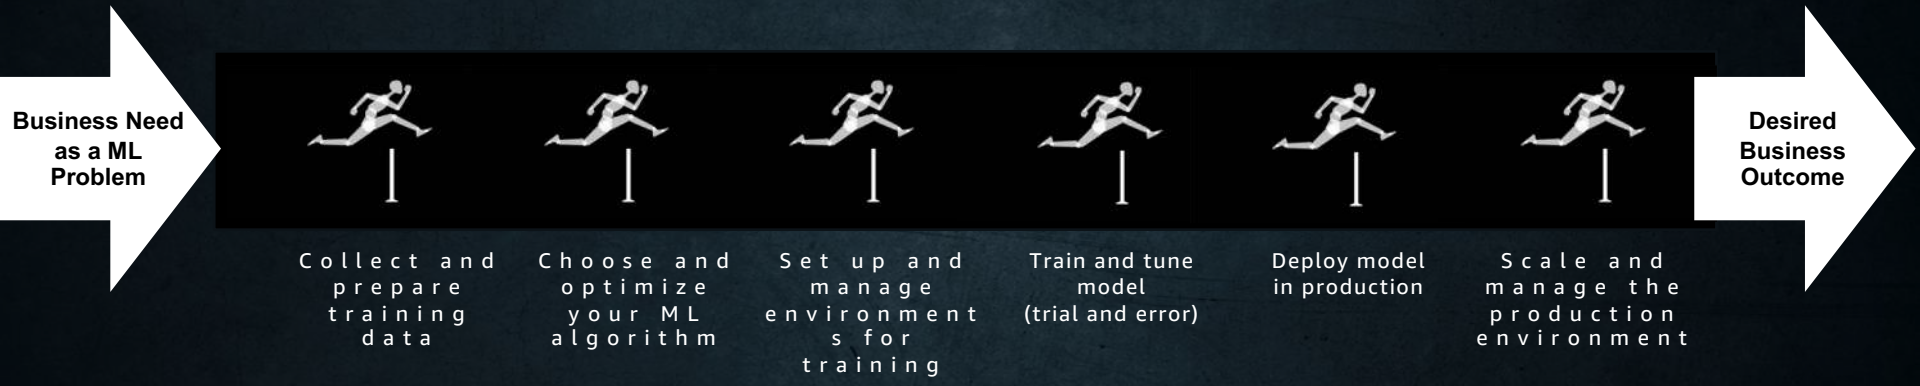

SIMPLIFY THE END-TO-END MACHINE LEARNING PROCESS

aws

Lab 1: Managing a high **quality** training set at **scale** using **SageMaker GroundTruth**

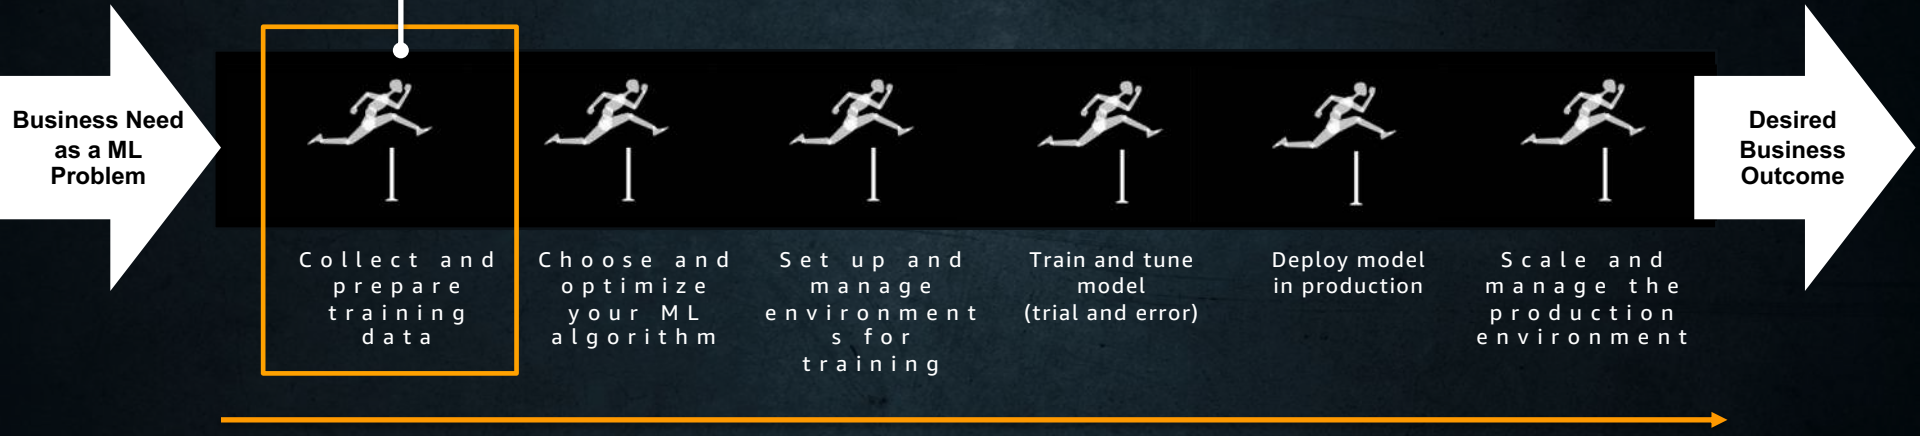

SIMPLIFY THE END-TO-END MACHINE LEARNING PROCESS

aws

Lab 2: Train, tune and deploy a custom object detector (**SSD**) with **zero coding**.

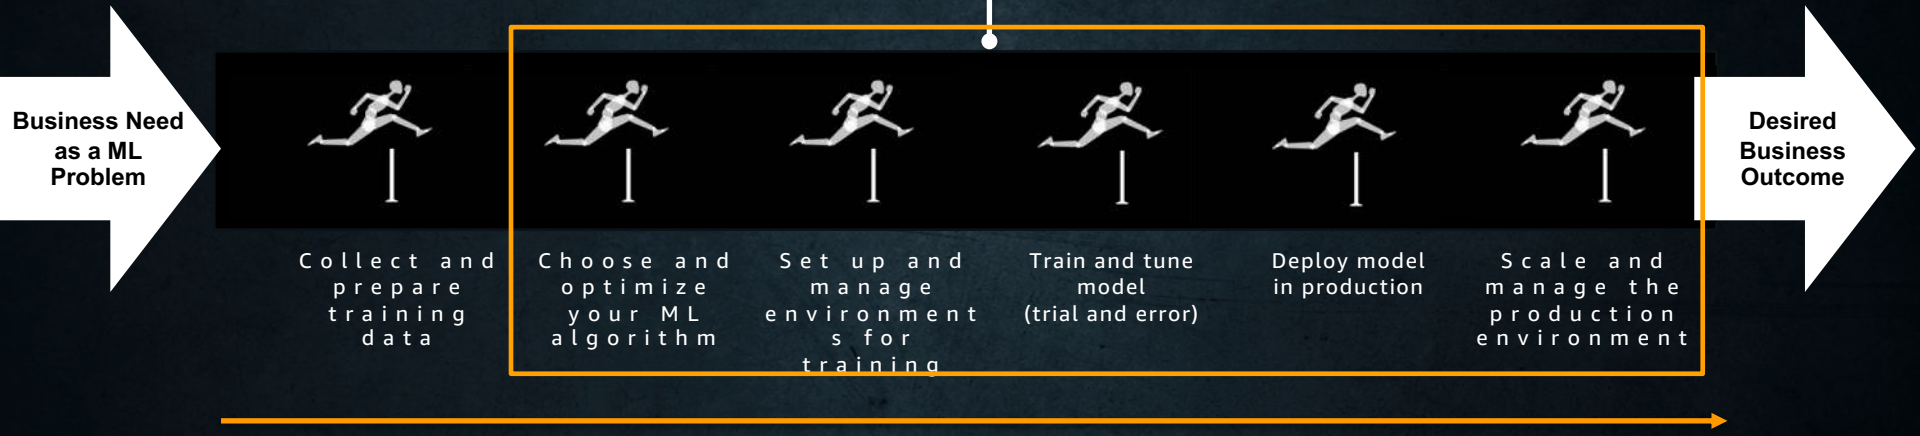

SIMPLIFY THE END-TO-END MACHINE LEARNING PROCESS

aws

Lab 3: "Bring Your Own Script:" train, tune and deploy a custom object detector (**YOLOv3**) on GluonCV (MXNet).

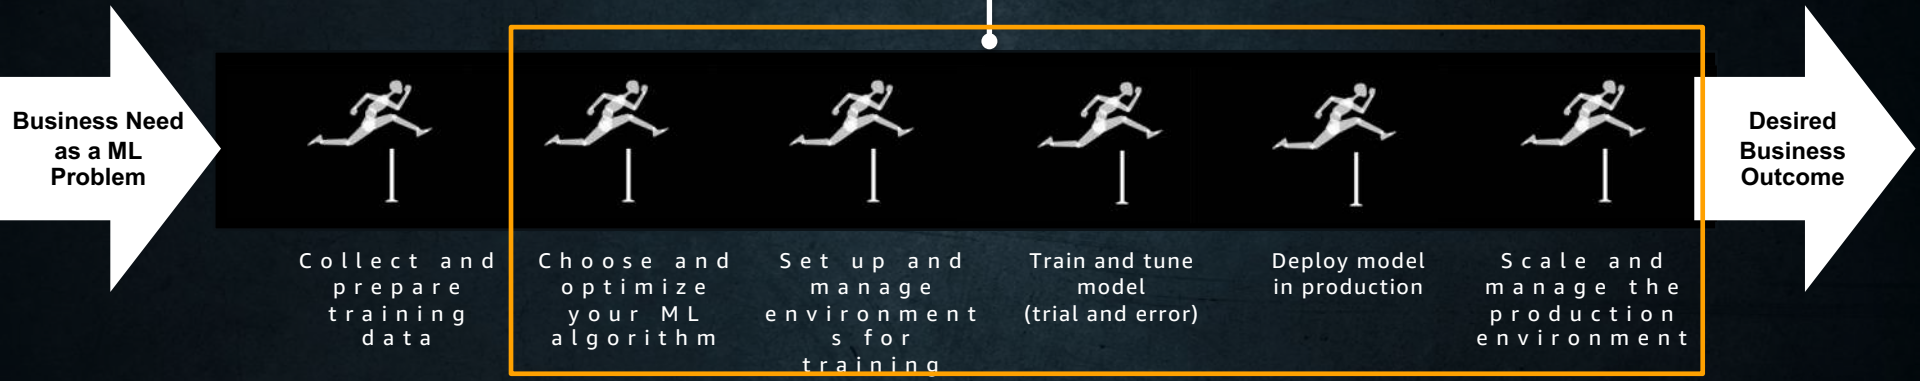

SIMPLIFY THE END-TO-END MACHINE LEARNING PROCESS

aws

GLUON

Supported Applications

Application	Illustration	Available Models
Image Classification: recognize an object in an image.		50+ models, including ResNet, MobileNet, DenseNet, VGG, ...
Object Detection: detect multiple objects with their bounding boxes in an image.	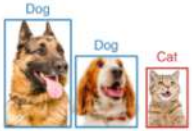	Faster RCNN, SSD, Yolo-v3
Semantic Segmentation: associate each pixel of an image with a categorical label.	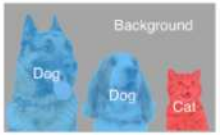	FCN, PSP, DeepLab v3
Instance Segmentation: associate each pixel of an image with an instance label.	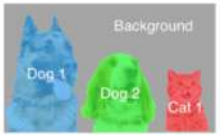	Mask RCNN
Pose Estimation: detect human pose from images.		Simple Pose

Algorithm variants : For instance, in **Object Detection**, different algorithms offer trade-offs between **accuracy (mAP)** and **latency (fps)**

aws

Future labs

- By feedback and demand! dylatong@amazon.com
- Recognition:
  - Face Search
  - People Tracking
- Textract solutions
- More Amazon SageMaker use cases:
  - Segmentation, Pose Estimation, Similarity Search

aws

Lab 1

Manage a high-quality data set at scale

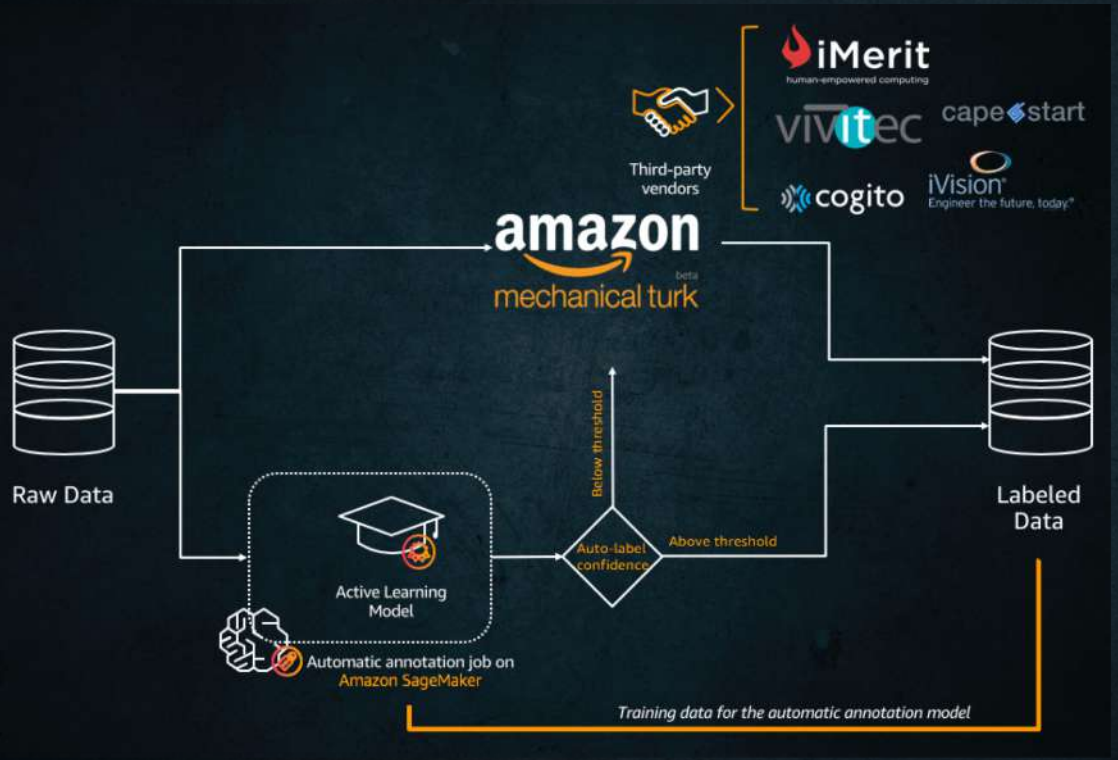

1. Launch a Notebook Instance.
2. Manage a private workforce.
3. Create an annotation job for Object Detection.
4. Generate a dataset and metadata compatible with Amazon SageMaker algorithms without further data wrangling!

aws

Lab 2

Create a custom object detector with zero coding

SAGEMAKER OPTIMIZED ALGORITHMS

- BlazingText Algorithm
- DeepAR Forecasting Algorithm
- Factorization Machines Algorithm
- Image Classification Algorithm
- IP Insights Algorithm
- K-Means Algorithm
- K-Nearest Neighbors (k-NN) Algorithm
- Latent Dirichlet Allocation (LDA) Algorithm
- Linear Learner Algorithm
- Neural Topic Model (NTM) Algorithm
- Object2Vec Algorithm
- **Object Detection Algorithm**
- Principal Component Analysis (PCA) Algorithm
- Random Cut Forest (RCF) Algorithm
- Semantic Segmentation Algorithm
- Sequence-to-Sequence Algorithm
- XGBoost Algorithm

1. Configure a hyperparameter tuning job for an Object Detection Algorithm.
2. Train on GPU
3. Deploy a managed endpoint.
4. Test and visualize!

aws

Lab 3

Bring Your Own Script and automate the ML process

Examples

Training: Only modifications required is to set script certain parameters values from SageMaker container environment variables.

GluonCV YOLOv3 training script
PyTorch Siamese Network training script

Inference: requires overriding programmatic interface implementation.

GluonCV YOLOv3 model serving script
PyTorch Siamese Network model serving script

- **input_fn**: request format pre-processing
- **model_fn:** how to load the model
- **predict_fn:** inference logic
- **output_fn:** response format processing

1. Bring your own YoloV3 script on GluonCV
2. Prepare your data set and environment
3. Explore and prototype locally
4. Automate model tuning, and train
5. Deploy, test and visualize!

aws

https://github.com/dylan-tong-aws/aws-cv-jumpstarter

aws.ai

Appendix

How we can help...

ML Solutions Lab

Brainstorming
Custom modeling
Training
Work side-by-side with Amazon experts

Machine Learning Training & Certification

Practical education on ML for new & experienced practitioners
Based on the same material used to train Amazon developers

aws